

Nicholas De Santos

DSC 680

Final Project 1

Background/History and our Business Problem

Great films have always influenced society. And vice versa. Films tell us a lot about the time in which they were released. You see things like Godzilla which on surface level is just a giant lizard terrorizing a city but in behind the scenes you see a comment on the unfortunate horrors of the atomic bomb attack against Japan. The event in turn influenced a lot of Japanese film. Especially horror. This just goes to show the power of movies. More importantly, the power of understanding how audiences feel about certain movies. Understanding what types of movies tend to be popular among the public also shines a light on what interests people the most. The purpose of this project is to investigate what sort of attributes and features a major film could have that would boost the film's ratings from its viewers.

Data Explanation (Data Prep, Data Dictionary, etc)

The data that we will be using a data set found on Kaggle.com titled "IMDb 5000+ Movies & Multiple Genres Dataset". This dataset has over 5000 different movie titles all containing information on the movie title, release year, director, actors, ratings, runtime in minutes, censor rating, total gross revenue, main genre and side genres (G, 2022). Below are the descriptions for each of the variables in the dataset.

'Movie_Title' : Consist of 5000+ Movie Titles (5000+ Unique Values)
'Year' : Ranging from 1920s to 2022 (99 Unique Values)
'Director' : Names the Director (2000+ Unique Values)
'Actors' : Names the Actors (5000+ Unique & Multiple Values)
'Rating' : Titles rated for 10 by 25k+ Voters (74+ Unique Values)
'main_genre' : Main Genre of the Title (13+ Unique Values)
'side_genre' : Side / Multiple Genre of the Movie (144+ Unique & Multiple Values)
'Runtime(Mins)' : Total duration of the movie in Minutes (156+ Unique Values)
'Censor' : Censorship of the Movie (25+ Unique Values)

'Total_Gross' : Total Box-Office Collection of the Movie (3500+ Unique Values)

Since the main goal of our investigation is to see what attributes and features are associated with more popular films, we will use the user rating variable "Rating" as our target variable when we get into our model building process.

Importing Data and Main Packages

Before we can begin to build our model we must first load our data. A view of what the data looks like can be found in the appendix at the end of this report. Our data will be saved into a dataframe object 'movies'.

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

#Loading Data
movies = pd.read_csv('movies.csv')
movies = pd.DataFrame(movies)
original_data = movies
```

Examining and Cleaning Data

Once our data has been loaded and saved in a dataframe, we can start examining the data more closely. Then we can see what we need to do in order to clean and prepare the data for our analysis.

The first thing that we did was we noticed that the total gross revenue variable was not a numeric variable so we took just the numbers and then turned the column into a numeric object. We can then check that the object conversion was successful as well.

```
In [2]: #Convert "Total_Gross" value to numerical variable
import re
movies['Total_Gross'] = movies['Total_Gross'].str.extract(r'(\d+\.\d+)')
movies['Total_Gross'] = pd.to_numeric(movies['Total_Gross'])
print(movies['Total_Gross'].dtype)

float64
```

In this next portion of the code we want to see all the different unique values for each of the categorical variables. This helps us decide to exclude the director variable later on in our model building process due to too many unique values. It no longer served as a category and was deemed unnecessary in the model building process as well. It

can also be useful in pointing out errors/misspelled strings in the categorical data.

```
In [3]: #Unique Values for Director
print(movies['Director'].value_counts())
```

```
Director
Woody Allen                40
Clint Eastwood             32
Steven Spielberg           28
Ron Howard                 24
Steven Soderbergh          23
..
Directors:Mike Judge, Mike de Seve, Brian Mulroney, Yvette Kaplan  1
Directors:Aaron Blaise, Robert Walker                             1
Stephen J. Anderson                                                1
Directors:Byron Howard, Chris Williams                           1
Gökhan Gök                                                         1
Name: count, Length: 2320, dtype: int64
```

```
In [4]: #Unique Values for Censor Rating
print(movies['Censor'].value_counts())
```

```
Censor
UA          1118
A           1101
U           1023
R            926
Not Rated   495
PG-13       405
18          136
PG          120
16          71
13          53
UA 16+       22
15+          18
7            17
UA 13+       12
G             9
(Banned)      8
UA 7+         7
12+           5
All           5
Unrated       4
U/A           2
18+           2
12            1
M/PG          1
NC-17         1
Name: count, dtype: int64
```

```
In [5]: #Unique Values for Main Genre
print(movies['main_genre'].value_counts())
```

```

main_genre
Action      1577
Comedy      1350
Drama       1027
Crime        447
Biography   355
Animation   321
Adventure   296
Horror       142
Mystery      26
Fantasy      13
Western       4
Film-Noir    3
Musical       1
Name: count, dtype: int64

```

In this portion of the code we are able to remove any duplicate rows in our dataset, remove any missing values within the dataset and removed both the actor, side genre and director our of our dataset. We also do some renaming of the columns within the dataset to make it easier to read and use within the code. We can finally have a look at what the final dataset looks like and move onto our data exploration and analysis.

```

In [6]: #Getting Rid of Duplicates: keeping first duplicate row
        movies = movies.drop_duplicates()

        #Getting Rid of Missing Value Rows: Based on NA values
        movies = movies.dropna()

        #Getting Rid of Actor and Side Genre columns
        movies = movies.drop(columns=['side_genre'])
        movies = movies.drop(columns=['Actors'])

        #Renaming Columns
        movies.rename(columns = {'Runtime(Mins)': 'Runtime'}, inplace = True)
        movies.rename(columns = {'Movie_Title': 'Title'}, inplace = True)
        movies.rename(columns = {'Total_Gross': 'GrossRev'}, inplace = True)
        movies.rename(columns = {'main_genre': 'Genre'}, inplace = True)

        movies

```

Out[6]:

	Title	Year	Director	Rating	Runtime	Censor	GrossRev	Genre
1	The Dark Knight	2008	Christopher Nolan	9.0	152	UA	534.86	Action
2	The Lord of the Rings: The Return of the King	2003	Peter Jackson	9.0	201	U	377.85	Action
3	Inception	2010	Christopher Nolan	8.8	148	UA	292.58	Action
4	The Lord of the Rings: The Two Towers	2002	Peter Jackson	8.8	179	UA	342.55	Action
5	The Lord of the Rings: The Fellowship of the Ring	2001	Peter Jackson	8.8	178	U	315.54	Action
...
5555	Son of the Mask	2005	Lawrence Guterman	2.2	94	U	17.02	Comedy
5557	Disaster Movie	2008	Directors:Jason Friedberg, Aaron Seltzer	1.9	87	PG-13	14.19	Comedy
5558	The Hottie & the Nottie	2008	Tom Putnam	1.9	91	PG-13	0.03	Comedy
5559	From Justin to Kelly	2003	Robert Iscove	1.9	81	PG	4.92	Comedy
5560	Superbabies: Baby Geniuses 2	2004	Bob Clark	1.5	88	PG	9.11	Comedy

4694 rows × 8 columns

Methods

Before getting into the model building process, we first conducted some exploratory data analysis on our data to investigate what kind of patterns and relationships we can find in the data before we create a model. It is in this state that we explored each variable's relationship with our target variable, including categorical variables. The purpose of those steps was to be able to give us a hint as to what variables should or shouldn't be included in the model building process as well as to shine a light on what variables are most important in the model.

After our data exploration, we finally get into our model creation process. Before we could use our data to train a model we first had to convert all of our categorical variables to dummy variable in order to be used for linear regression in the model building process.

We then split that new dataframe with dummy variables instead of

categorical variables into training and testing sets of the original data. We built our model with the training dataset and then later evaluated that model using the RMSE value of our testing set compared to the model predictions.

Analysis

Data Exploration

Next we move onto exploring any possible relationships between the variables. To do so we look at both the scatter plot graph of each variable along with our target variable "Rating" along with the correlation coefficient between those two variables.

Numerical Variables

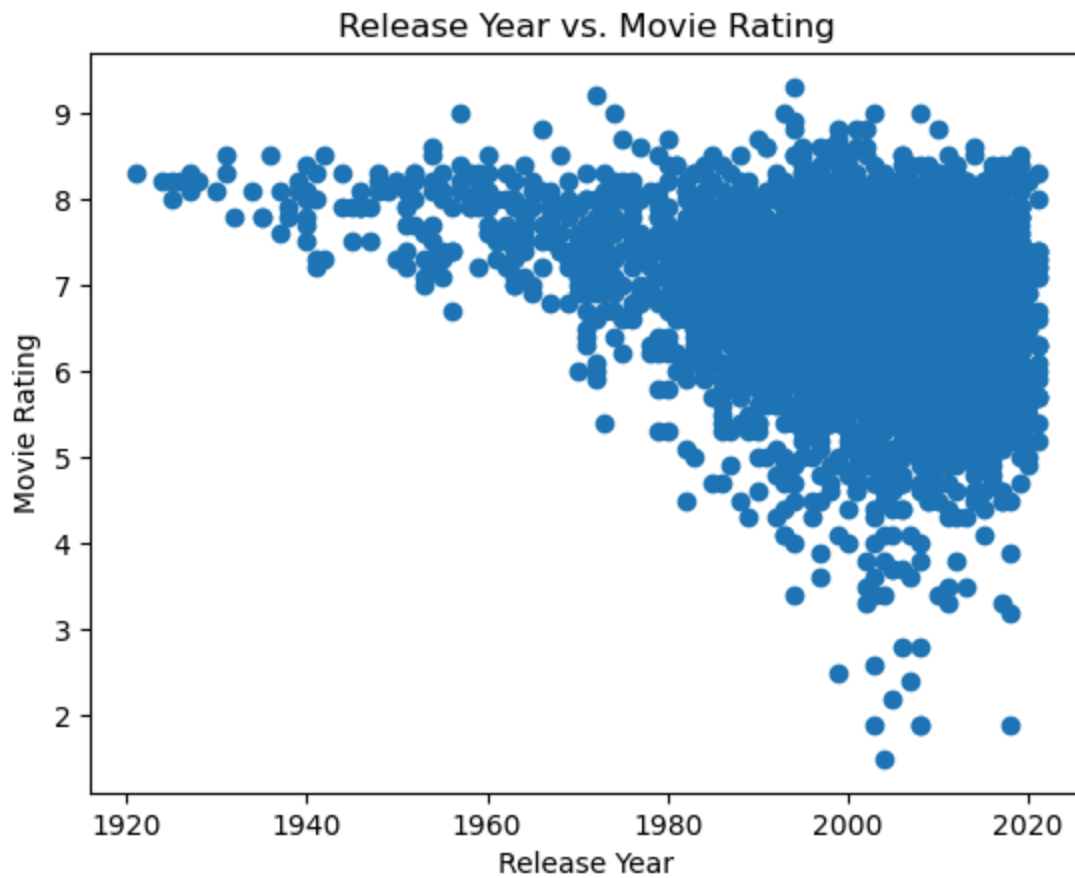
```
In [7]: #Exploring Relationships: Year vs Rating

fig, ax = plt.subplots()
plt.scatter(movies.Year, movies.Rating)
plt.xlabel("Release Year")
plt.ylabel("Movie Rating")
plt.title("Release Year vs. Movie Rating")

plt.show()

import scipy.stats as stats

#correlation coefficient and PValue
correlation = stats.pearsonr(movies.Year, movies.Rating)
print(correlation)
```



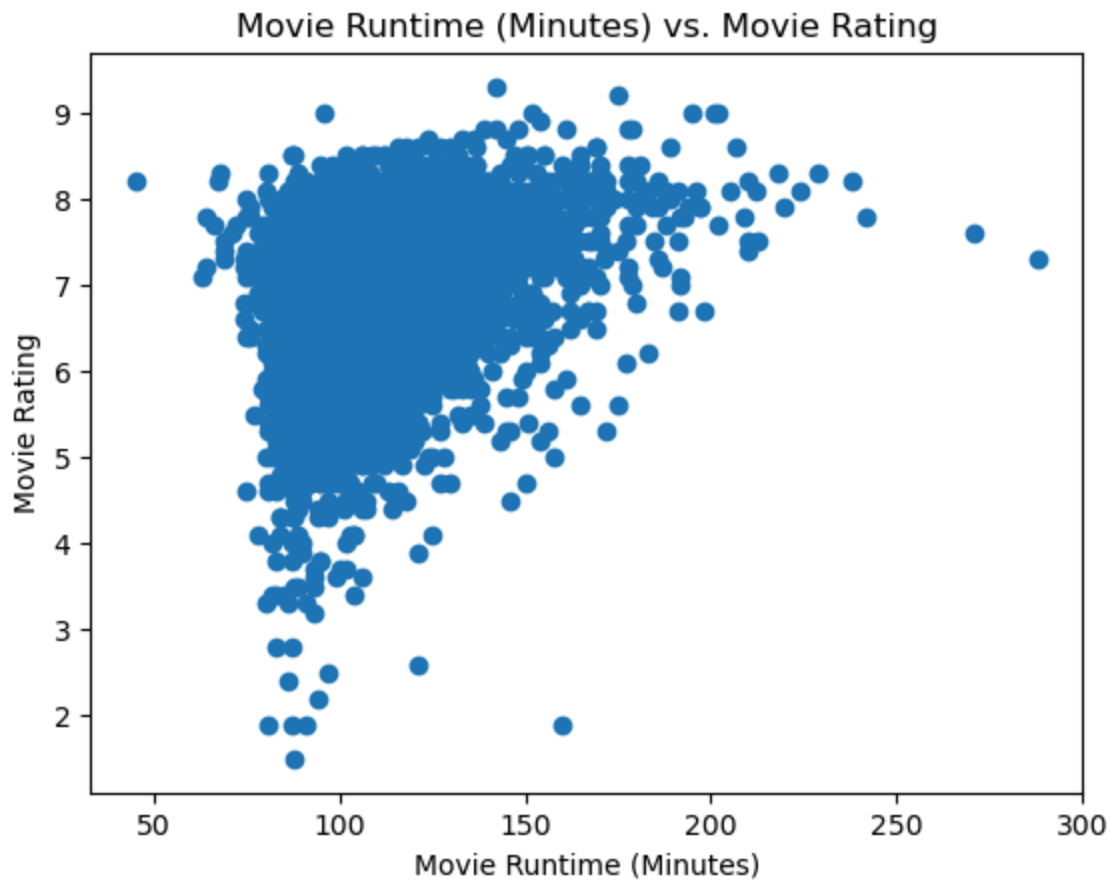
PearsonRResult(statistic=-0.2631474562741687, pvalue=3.3952118104659924e-75)

```
In [8]: #Exploring Relationships: Runtime vs Rating

fig, ax = plt.subplots()
plt.scatter(movies.Runtime, movies.Rating)
plt.xlabel("Movie Runtime (Minutes)")
plt.ylabel("Movie Rating")
plt.title("Movie Runtime (Minutes) vs. Movie Rating")

plt.show()

#correlation coefficient and PValue
correlation = stats.pearsonr(movies.Runtime, movies.Rating)
print(correlation)
```



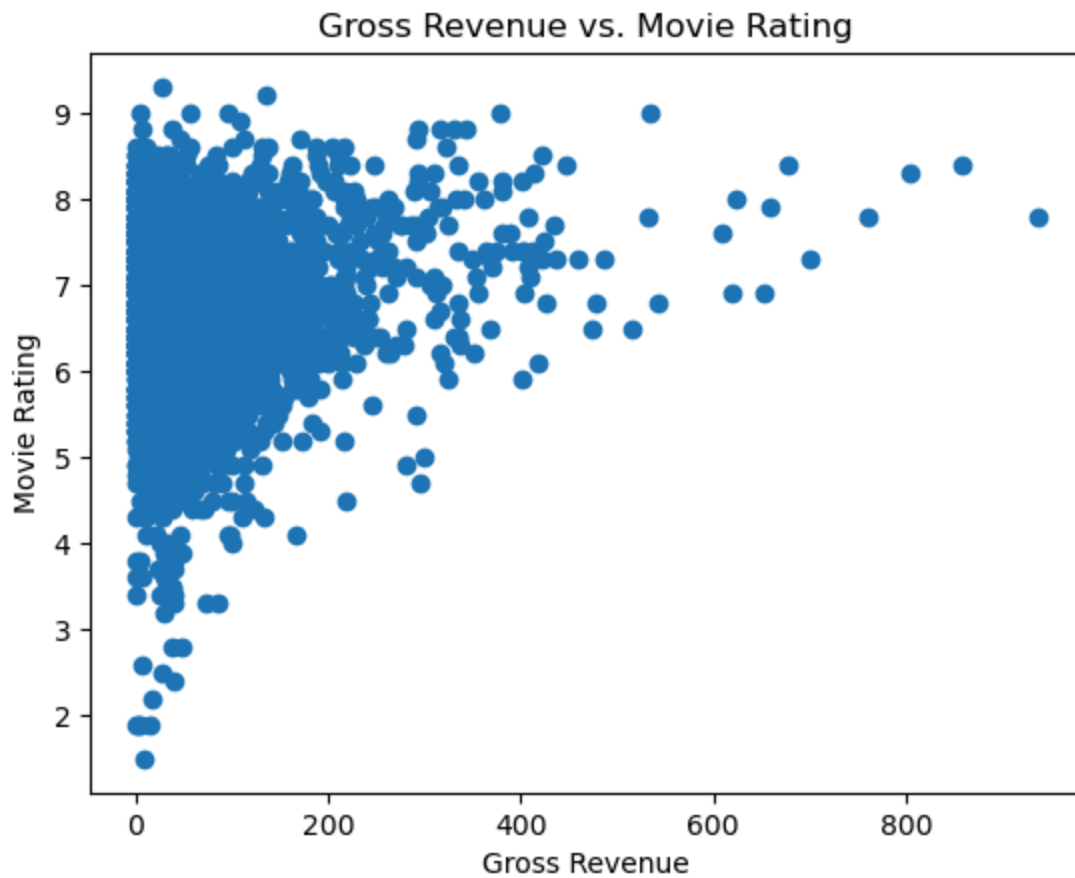
PearsonRResult(statistic=0.35976408087247, pvalue=1.944673821726563e-143)

```
In [9]: #Exploring Relationships: Year vs Rating

fig, ax = plt.subplots()
plt.scatter(movies.GrossRev, movies.Rating)
plt.xlabel("Gross Revenue")
plt.ylabel("Movie Rating")
plt.title("Gross Revenue vs. Movie Rating")

plt.show()

#correlation coefficient and PValue
correlation = stats.pearsonr(movies.GrossRev, movies.Rating)
print(correlation)
```

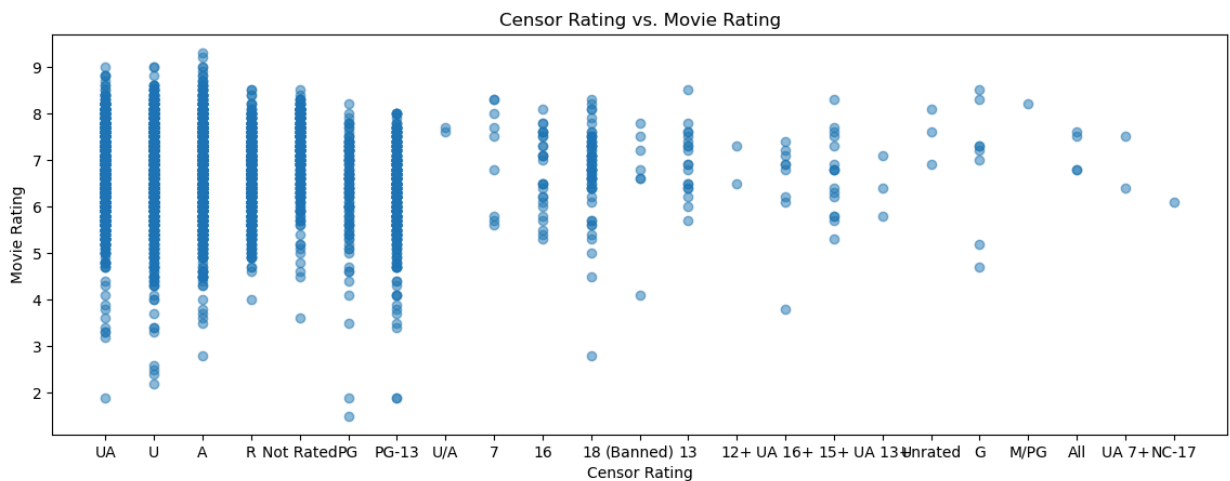
PearsonRResult(statistic=0.07276128277732986, pvalue=6.026534178753546e-07)

Categorical Variables

In [13]: *#Exploring Relationships: Censor Rating vs Movie Rating*

```
fig, ax = plt.subplots()
fig.set_figwidth(14)
plt.scatter(movies.Censor, movies.Rating, alpha = 0.5)
plt.xlabel("Censor Rating")
plt.ylabel("Movie Rating")
plt.title("Censor Rating vs. Movie Rating")

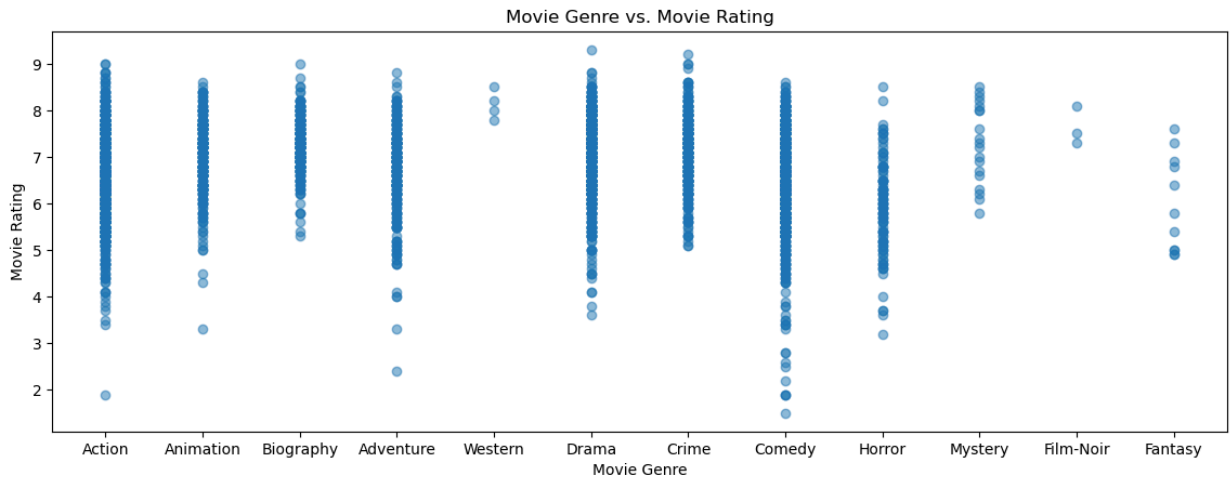
plt.show()
```



In [12]: *#Exploring Relationships: Genre vs Movie Rating*

```
fig, ax = plt.subplots()
fig.set_figwidth(14)
plt.scatter(movies.Genre, movies.Rating, alpha = 0.5)
plt.xlabel("Movie Genre")
plt.ylabel("Movie Rating")
plt.title("Movie Genre vs. Movie Rating")

plt.show()
```



Model Creation and Testing

Finally, we get to our model building process in order to try and predict the viewer rating of a film. Before we start the process we first want to get the correct formatting and object types for model creation. We first start with turning our categorical variables into dummy variables that we can use for linear regression.

In [14]: *#Turn categorical variables to dummy variables for model creation process*

```
moviemod = movies
from sklearn.preprocessing import LabelBinarizer

label_binarizer = LabelBinarizer()
label_binarizer_output = label_binarizer.fit_transform(moviemod['Censor'])
censord = pd.DataFrame(label_binarizer_output, columns = label_binarizer.classes_)

label_binarizer_output = label_binarizer.fit_transform(moviemod['Genre'])
genred = pd.DataFrame(label_binarizer_output, columns = label_binarizer.classes_)

moviemod = moviemod.drop(['Title', 'Director', 'Censor', 'Genre'], axis=1)
moviemod = pd.concat([moviemod, censord, genred], axis=1, join='inner')
moviemod
```

Out[14]:

	Year	Rating	Runtime	GrossRev	(Banned)	12+	13	15+	16	18	...	Animation	Biography
1	2008	9.0	152	534.86	0	0	0	0	0	0	...	0	0
2	2003	9.0	201	377.85	0	0	0	0	0	0	...	0	0
3	2010	8.8	148	292.58	0	0	0	0	0	0	...	0	0
4	2002	8.8	179	342.55	0	0	0	0	0	0	...	0	0
5	2001	8.8	178	315.54	0	0	0	0	0	0	...	0	0
...
4689	2003	6.1	92	36.92	0	0	0	0	0	0	...	0	0
4690	2002	6.1	117	9.68	0	0	0	0	0	0	...	0	0
4691	2001	6.1	109	13.73	0	0	0	0	0	0	...	0	0
4692	2000	6.1	103	37.17	0	0	0	0	0	0	...	0	0
4693	1999	6.1	122	22.36	0	0	0	0	0	0	...	0	0

3936 rows × 39 columns

Next we split our data into two different sets. The first set is the training set that we will use to make our model. The other set is the testing set. The testing set is used to test the model's performance.

```
In [15]: #Splitting Training and Testing Sets

from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split

# get the locations
X = moviemod.drop(['Rating'], axis=1)
y = moviemod['Rating']

# split the dataset
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.05, random_state=42)
```

Results and Conclusion

```
In [16]: #Creating Linear Regression Model

import sklearn.metrics as metrics
from sklearn.linear_model import LinearRegression

model = LinearRegression()
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
```

```
# RMSE  
print(np.sqrt(metrics.mean_squared_error(y_test, y_pred)))
```

0.6971517528316209

As you can see, we used our training x variables of our data to create our model which later produced viewer rating predictions for the various attributes in our testing x variables. We then used those predictions and compared them to the actual target variable values (y_test) and got an RMSE value of about 0.6972. The cut off of what is an "acceptable" RMSE value is typically 0.5 so you can see that we were close but not quite there. This might be a good reason to future the investigation and work on an improved model.

Assumptions

For our data we assume that the information about the films listed in the dataset are completely accurate and up to date as far as the documentation states in the Kaggle description. We also assume all the necessary assumption to use linear regression in our model building process. The assumptions include the assumption that the chosen sample is representative of the population. The assumption that there is a linear relationship between the independent variable(s) and the dependent variable. And the assumption that all the variables are normally distributed.

Limitations

Some limitations to this project and investigation was the data that we were able to use, the time given to hold the investigation and get our results, and our technical skills with Python as a programming language.

Challenges

The main challenges that we faced was mostly the data preparation and cleaning. It was difficult deciding what columns/variables we would keep in our model as well.

Future Uses/Additional Applications

In terms of what this model could be used for, the model we created can be used to predict the popularity of new films based on given features which was one of our original goals.

The model can also be used to compare films actual performance to what the model predicted and use those results to continue to better our model and increase the accuracy.

Ethical Assessment

In terms of how ethical this process was. I don't believe there was much of a breach in private information. we believe because the information provided for all the movie titles can be easily found online there are no ethical concerns for the names of actors or directors listed in the dataset. There could be some copyright and property issues if this model were to be used in some kind of business advancement in the film industry without addressing the rightful owners that intellectual property. That is a valid concern.

References

G, R. A. (2022, October 29). IMDB 5000+ movies & multiple genres dataset. Kaggle.
<https://www.kaggle.com/datasets/rakkesharv/imdb-5000-movies-multiple-genres-dataset?resource=download>

Appendix

Table 1: Original Dataframe

In [23]: `original_data`

Out[23]:

	Movie_Title	Year	Director	Actors	Rating	Runtime(Mins)	Censor	Total_Gross	mai
0	Kantara	2022	Rishab Shetty	Rishab Shetty, Sapthami Gowda, Kishore Kumar G...	9.3	148	UA	NaN	
1	The Dark Knight	2008	Christopher Nolan	Christian Bale, Heath Ledger, Aaron Eckhart, M...	9.0	152	UA	534.86	
2	The Lord of the Rings: The Return of the King	2003	Peter Jackson	Elijah Wood, Viggo Mortensen, Ian McKellen, Or...	9.0	201	U	377.85	
3	Inception	2010	Christopher Nolan	Leonardo DiCaprio, Joseph Gordon-Levitt, Ellio...	8.8	148	UA	292.58	
4	The Lord of the Rings: The Two Towers	2002	Peter Jackson	Elijah Wood, Ian McKellen, Viggo Mortensen, Or...	8.8	179	UA	342.55	
...
5557	Disaster Movie	2008	Directors:Jason Friedberg, Aaron Seltzer	Carmen Electra, Vanessa Lachey, Nicole Parker,...	1.9	87	PG-13	14.19	
5558	The Hottie & the Nottie	2008	Tom Putnam	Paris Hilton, Joel David Moore, Christine Laki...	1.9	91	PG-13	0.03	
5559	From Justin to Kelly	2003	Robert Iscove	Kelly Clarkson, Justin Guarini, Katherine Bail...	1.9	81	PG	4.92	
5560	Superbabies: Baby	2004	Bob Clark	Jon Voight, Scott Baio,	1.5	88	PG	9.11	

	Movie_Title	Year	Director	Actors	Rating	Runtime(Mins)	Censor	Total_Gross	mai
	Geniuses 2			Vanessa Angel, Skyler ...					
5561	Cumali Ceber: Allah Seni Alsin	2017	Gökhan Gök	Halil Söyletmez, Doga Konakoglu, Emre Keskin, ...	1.0	100	Not Rated	NaN	
5562	10	1							